

# Machine Learning Essentials SS25 - Exercise Sheet 6

## Instructions

- TODO 's indicate where you need to complete the implementations.
- You may use external resources, but **write your own solutions**.
- Provide concise, but comprehensible comments to explain what your code does.
- Code that's unnecessarily extensive and/or not well commented will not be scored.

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import torch as tc
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from torchsummary import summary

np.random.seed(42)
tc.manual_seed(42)

device = tc.device("cuda" if tc.cuda.is_available() else "cpu")
```

## Exercise 2 - CNN Classifier

The SIGNS dataset is a collection of 6 signs representing numbers from 0 to 5. We first load the data and have the shapes printed out. The split into train, validation and test set has already been carried out.

```

In [2]: # Load the dataset
X_train = np.load('sign_data/X_train.npy')
Y_train = np.load('sign_data/Y_train.npy')
X_val = np.load('sign_data/X_val.npy')
Y_val = np.load('sign_data/Y_val.npy')
X_test = np.load('sign_data/X_test.npy')
Y_test = np.load('sign_data/Y_test.npy')

# print the shape of the dataset
print("X_train shape: " + str(X_train.shape))
print("Y_train shape: " + str(Y_train.shape))
print("X_val shape: " + str(X_val.shape))
print("Y_val shape: " + str(Y_val.shape))
print("X_test shape: " + str(X_test.shape))
print("Y_test shape: " + str(Y_test.shape)+"\n")
print("classes: " + str(np.unique(Y_train)))

# check if classes are balanced
print("Counts of classes in Y_train: " + str(np.unique(Y_train, return_co
print("Counts of classes in Y_val: " + str(np.unique(Y_val, return_counts
print("Counts of classes in Y_test: " + str(np.unique(Y_test, return_coun

```

X\_train shape: (960, 64, 64, 3)

Y\_train shape: (960,)

X\_val shape: (120, 64, 64, 3)

Y\_val shape: (120,)

X\_test shape: (120, 64, 64, 3)

Y\_test shape: (120,)

classes: [0 1 2 3 4 5]

Counts of classes in Y\_train: [160 160 160 160 160 160]

Counts of classes in Y\_val: [20 20 20 20 20 20]


Counts of classes in Y\_test: [20 20 20 20 20 20]

The classes are balanced so that accuracy is an appropriate measure for evaluating a classifier. We next visualize an instance of each class.

```

In [3]: fig, axs = plt.subplots(1, 6, figsize=(20, 10))
for i in range(6):
    # get indices where the label is i
    idx = np.where(Y_train == i)[0][0]
    axs[i].imshow(X_train[idx])
    axs[i].set_title("y = " + str(i))

```



Pixels in each channel (RGB) of the images take values in the range  $[0, 255]$ . However, it is desirable to have absolute values in the range  $[0, 1]$  as input for neural network architectures to avoid exploding or vanishing gradient problems. Through the following cell, we apply a simple data scaling procedure: we divide the values of the pixels by 255. As an alternative, you can use the `StandardScaler()` function of the scikit-learn library.

```
In [4]: X_train = X_train/255  
X_val = X_val/255  
X_test = X_test/255
```

```
In [5]: X_train[0].shape
```

```
Out[5]: (64, 64, 3)
```

## Task 1

Use pytorch to build the model. Take a look at the [documentation](#) for an introduction, a detailed tutorial, for example for classifiers, can be found [here](#).

Implement the following architecture:

- Conv2d: 4 output channels, 3 by 3 filter size, stride 1, padding "same"
- BatchNorm2d: 4 output channels
- ReLU activation
- MaxPool2d: 2 by 2 filter size, stride 2, padding 0
- Conv2d: 8 output channels, 3 by 3 filter size, stride 1, padding "same"
- BatchNorm2d: 8 output channels
- ReLU activation
- MaxPool2d: Use a 2 by 2 filter size, stride 2, padding 0
- Flatten the previous output
- Linear: 64 output neurons
- ReLU activation function
- Linear: 6 output neurons
- LogSoftmax

We use the [LogSoftmax](#) here instead of the Softmax for computational reasons. Accordingly, the loss function is not CrossEntropyLoss but NLLLoss. When flattening, be careful not to do it with the batch dimension but only with the height, width and channel dimension.

```
In [6]: class CNN_Classifier(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.first_conv = nn.Sequential(  
            nn.Conv2d(3, 4, (3, 3), 1, "same"),  
            nn.BatchNorm2d(4),  
            nn.ReLU(),  
            nn.MaxPool2d(2, 2, 0),  
        )  
        # output dim 32 [(64-2)/2 + 1]  
  
        self.second_conv = nn.Sequential(  
            nn.Conv2d(4, 8, (3,3), 1, "same"),  
            nn.BatchNorm2d(8),  
            nn.ReLU(),  
            nn.MaxPool2d(2, 2, 0),  
        )
```

```

    )
    # output dim = 16 [(32 - 2)/2 + 1]

    self.lin_layer = nn.Sequential(
        nn.Flatten(),
        # output dim = 16 * 16 * 8 = 2048
        nn.Linear(2048, 64),
        nn.ReLU(),
        nn.Linear(64, 6)
    )

    def forward(self, X):
        # TODO: Implement the forward pass
        X = self.first_conv(X)
        X = self.second_conv(X)
        X = self.lin_layer(X)
        return nn.LogSoftmax(dim=1)(X)

```

To test your model you can forward some random numbers. The shape of the output should be (2, 6).

```

In [7]: cnn_model = CNN_Classifier()
        # dummy sample of batch size 2
        X_random = tc.randn(2, 3, 64, 64)
        output = cnn_model(X_random)

        print("Output shape: " + str(output.shape))

```

Output shape: torch.Size([2, 6])

torchsummary.summary provides a nice overview of the model and the number of learnable parameters:

```

In [8]: summary(cnn_model, input_size=(3, 64, 64), device="cpu")

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 4, 64, 64]	112
BatchNorm2d-2	[-1, 4, 64, 64]	8
ReLU-3	[-1, 4, 64, 64]	0
MaxPool2d-4	[-1, 4, 32, 32]	0
Conv2d-5	[-1, 8, 32, 32]	296
BatchNorm2d-6	[-1, 8, 32, 32]	16
ReLU-7	[-1, 8, 32, 32]	0
MaxPool2d-8	[-1, 8, 16, 16]	0
Flatten-9	[-1, 2048]	0
Linear-10	[-1, 64]	131,136
ReLU-11	[-1, 64]	0
Linear-12	[-1, 6]	390

Total params: 131,958  
 Trainable params: 131,958  
 Non-trainable params: 0

Input size (MB): 0.05  
 Forward/backward pass size (MB): 0.63  
 Params size (MB): 0.50  
 Estimated Total Size (MB): 1.18

## Task 2

DataLoaders wrap around Datasets to provide efficient data batching, shuffling, and parallel loading during model training or inference. To define a custom dataset we must implement three functions: **init**, **len** and **get\_item**. While **len** defines the length of the dataset and thus the number of batches in the dataloader, **get\_item** can be used to get a single sample through the index.

```
In [9]: class Image_Dataset(Dataset):
        def __init__(self, X, Y):
            self.X = X
            self.Y = Y

        def __len__(self):
            # TODO: Return the length of the dataset
            return len(self.X)

        def __getitem__(self, idx):
            # TODO: Return the image as a float tensor and the label as a long
            img = self.X[idx]
            label = np.array(self.Y[idx])
            img = tc.from_numpy(img).float()

            #print(label.size())

            # permute the shape to (3, 64, 64)
            img = img.permute(2, 0, 1)
            label = tc.from_numpy(label).long()
            return img, label
```

```
In [10]: train_batch_size = 64
```

```

val_batch_size = len(Y_val)
test_batch_size = len(Y_test)

# TODO: Create the dataset and dataloader
train_dataset = Image_Dataset(X_train, Y_train)
val_dataset = Image_Dataset(X_val, Y_val)
test_dataset = Image_Dataset(X_test, Y_test)
train_loader = DataLoader(train_dataset, train_batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, val_batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, test_batch_size, shuffle=True)

```

To make sure that everything has worked properly, we take a sample of the data\_loader and visualize it.


```

In [11]: sample_X, sample_Y = next(iter(train_loader))
plt.imshow(sample_X[0].T)
plt.title("y = " + str(int(sample_Y[0].item())))
plt.show()

```

/tmp/ipykernel\_1184/2461653407.py:2: UserWarning: The use of `x.T` on tensors of dimension other than 2 to reverse their shape is deprecated and it will throw an error in a future release. Consider `x.mT` to transpose batches of matrices or `x.permute(\*torch.arange(x.ndim - 1, -1, -1))` to reverse the dimensions of a tensor. (Triggered internally at /pytorch/aten/src/ATen/native/TensorShape.cpp:4413.)

```
plt.imshow(sample_X[0].T)
```



### Task 3

Implement the training loop. Use the negative log-likelihood loss (NLLLoss) and the Adam optimizer. Be sure to zero the gradients after each optimization step to avoid accumulating contributions from previous epochs and batches.

```
In [12]: def train_cnn(model, train_loader, val_loader, lr, n_epochs, device):
    model = model.to(device)

    # TODO: Initialize the optimizer and loss function
    loss_function = nn.NLLLoss()
    optimizer = optim.Adam(model.parameters(), lr=lr)

    train_loss = np.zeros(n_epochs)
    val_loss = np.zeros(n_epochs)
    train_acc = np.zeros(n_epochs)
    val_acc = np.zeros(n_epochs)

    for epoch in range(1, n_epochs + 1):
        model.train()

        epoch_loss = 0

        for X, Y in train_loader:
            X, Y = X.to(device), Y.to(device)
            # TODO: Implement the training loop

            optimizer.zero_grad()

            output = model(X)

            loss = loss_function(output, Y)

            loss.backward()

            optimizer.step()

            epoch_loss += loss.item()/len(train_loader)

        train_loss[epoch - 1] = epoch_loss
        train_acc[epoch - 1] = (output.argmax(dim=1) == Y).float().mean()

        model.eval()

        with tc.no_grad():
            X, Y = next(iter(val_loader))
            X, Y = X.to(device), Y.to(device)
            # TODO: Implement the evaluation step
            output = model(X)
            loss = loss_function(output, Y)

            val_loss[epoch - 1] = loss.item()
            val_acc[epoch - 1] = (output.argmax(dim=1) == Y).float().mean()
        print(f"Epoch {epoch}/{n_epochs} - Train Loss: {epoch_loss:.4f},

    return train_loss, val_loss, train_acc, val_acc
```

```
In [13]: n_epochs = 50
# TODO: Train the model with different learning rates
```

```
learning_rates = [0.00001, 0.0001, 0.001, 0.1]
train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []

models = [CNN_Classifier() for _ in learning_rates]
for lr in learning_rates:
    model = models[learning_rates.index(lr)]
    train_loss, val_loss, train_acc, val_acc = train_cnn(model, train_loader, val_loader)
    train_losses.append(train_loss)
    val_losses.append(val_loss)
    train_accuracies.append(train_acc)
    val_accuracies.append(val_acc)
```



Epoch 1/50 - Train Loss: 1.7900, Test Loss: 1.7928  
Epoch 2/50 - Train Loss: 1.7706, Test Loss: 1.7889  
Epoch 3/50 - Train Loss: 1.7515, Test Loss: 1.7732  
Epoch 4/50 - Train Loss: 1.7318, Test Loss: 1.7429  
Epoch 5/50 - Train Loss: 1.7109, Test Loss: 1.7155  
Epoch 6/50 - Train Loss: 1.6911, Test Loss: 1.6952  
Epoch 7/50 - Train Loss: 1.6717, Test Loss: 1.6775  
Epoch 8/50 - Train Loss: 1.6523, Test Loss: 1.6594  
Epoch 9/50 - Train Loss: 1.6329, Test Loss: 1.6422  
Epoch 10/50 - Train Loss: 1.6127, Test Loss: 1.6240  
Epoch 11/50 - Train Loss: 1.5940, Test Loss: 1.6060  
Epoch 12/50 - Train Loss: 1.5749, Test Loss: 1.5880  
Epoch 13/50 - Train Loss: 1.5563, Test Loss: 1.5711  
Epoch 14/50 - Train Loss: 1.5377, Test Loss: 1.5546  
Epoch 15/50 - Train Loss: 1.5185, Test Loss: 1.5382  
Epoch 16/50 - Train Loss: 1.5005, Test Loss: 1.5214  
Epoch 17/50 - Train Loss: 1.4825, Test Loss: 1.5052  
Epoch 18/50 - Train Loss: 1.4637, Test Loss: 1.4895  
Epoch 19/50 - Train Loss: 1.4476, Test Loss: 1.4744  
Epoch 20/50 - Train Loss: 1.4303, Test Loss: 1.4584  
Epoch 21/50 - Train Loss: 1.4126, Test Loss: 1.4437  
Epoch 22/50 - Train Loss: 1.3955, Test Loss: 1.4292  
Epoch 23/50 - Train Loss: 1.3794, Test Loss: 1.4140  
Epoch 24/50 - Train Loss: 1.3633, Test Loss: 1.3998  
Epoch 25/50 - Train Loss: 1.3459, Test Loss: 1.3849  
Epoch 26/50 - Train Loss: 1.3298, Test Loss: 1.3706  
Epoch 27/50 - Train Loss: 1.3131, Test Loss: 1.3565  
Epoch 28/50 - Train Loss: 1.2966, Test Loss: 1.3424  
Epoch 29/50 - Train Loss: 1.2812, Test Loss: 1.3286  
Epoch 30/50 - Train Loss: 1.2654, Test Loss: 1.3156  
Epoch 31/50 - Train Loss: 1.2515, Test Loss: 1.3024  
Epoch 32/50 - Train Loss: 1.2355, Test Loss: 1.2892  
Epoch 33/50 - Train Loss: 1.2210, Test Loss: 1.2759  
Epoch 34/50 - Train Loss: 1.2069, Test Loss: 1.2634  
Epoch 35/50 - Train Loss: 1.1923, Test Loss: 1.2503  
Epoch 36/50 - Train Loss: 1.1774, Test Loss: 1.2388  
Epoch 37/50 - Train Loss: 1.1633, Test Loss: 1.2263  
Epoch 38/50 - Train Loss: 1.1507, Test Loss: 1.2139  
Epoch 39/50 - Train Loss: 1.1362, Test Loss: 1.2013  
Epoch 40/50 - Train Loss: 1.1217, Test Loss: 1.1895  
Epoch 41/50 - Train Loss: 1.1086, Test Loss: 1.1784  
Epoch 42/50 - Train Loss: 1.0950, Test Loss: 1.1654  
Epoch 43/50 - Train Loss: 1.0829, Test Loss: 1.1542  
Epoch 44/50 - Train Loss: 1.0702, Test Loss: 1.1441  
Epoch 45/50 - Train Loss: 1.0581, Test Loss: 1.1335  
Epoch 46/50 - Train Loss: 1.0446, Test Loss: 1.1220  
Epoch 47/50 - Train Loss: 1.0339, Test Loss: 1.1125  
Epoch 48/50 - Train Loss: 1.0229, Test Loss: 1.1024  
Epoch 49/50 - Train Loss: 1.0103, Test Loss: 1.0937  
Epoch 50/50 - Train Loss: 0.9999, Test Loss: 1.0852  
Epoch 1/50 - Train Loss: 1.7328, Test Loss: 1.7763  
Epoch 2/50 - Train Loss: 1.5836, Test Loss: 1.6818  
Epoch 3/50 - Train Loss: 1.4290, Test Loss: 1.4960  
Epoch 4/50 - Train Loss: 1.2821, Test Loss: 1.3244  
Epoch 5/50 - Train Loss: 1.1452, Test Loss: 1.1932  
Epoch 6/50 - Train Loss: 1.0275, Test Loss: 1.0893  
Epoch 7/50 - Train Loss: 0.9219, Test Loss: 1.0050  
Epoch 8/50 - Train Loss: 0.8322, Test Loss: 0.9339  
Epoch 9/50 - Train Loss: 0.7531, Test Loss: 0.8691  
Epoch 10/50 - Train Loss: 0.6791, Test Loss: 0.8124

Epoch 11/50	- Train Loss: 0.6182, Test Loss: 0.7605
Epoch 12/50	- Train Loss: 0.5684, Test Loss: 0.7257
Epoch 13/50	- Train Loss: 0.5235, Test Loss: 0.6908
Epoch 14/50	- Train Loss: 0.4764, Test Loss: 0.6575
Epoch 15/50	- Train Loss: 0.4430, Test Loss: 0.6312
Epoch 16/50	- Train Loss: 0.4082, Test Loss: 0.6031
Epoch 17/50	- Train Loss: 0.3759, Test Loss: 0.5758
Epoch 18/50	- Train Loss: 0.3474, Test Loss: 0.5696
Epoch 19/50	- Train Loss: 0.3236, Test Loss: 0.5438
Epoch 20/50	- Train Loss: 0.3025, Test Loss: 0.5231
Epoch 21/50	- Train Loss: 0.2816, Test Loss: 0.5069
Epoch 22/50	- Train Loss: 0.2647, Test Loss: 0.4935
Epoch 23/50	- Train Loss: 0.2524, Test Loss: 0.4739
Epoch 24/50	- Train Loss: 0.2338, Test Loss: 0.4992
Epoch 25/50	- Train Loss: 0.2192, Test Loss: 0.4479
Epoch 26/50	- Train Loss: 0.2026, Test Loss: 0.4497
Epoch 27/50	- Train Loss: 0.1934, Test Loss: 0.4371
Epoch 28/50	- Train Loss: 0.1814, Test Loss: 0.4181
Epoch 29/50	- Train Loss: 0.1686, Test Loss: 0.4272
Epoch 30/50	- Train Loss: 0.1601, Test Loss: 0.4147
Epoch 31/50	- Train Loss: 0.1506, Test Loss: 0.3947
Epoch 32/50	- Train Loss: 0.1417, Test Loss: 0.4022
Epoch 33/50	- Train Loss: 0.1352, Test Loss: 0.3811
Epoch 34/50	- Train Loss: 0.1251, Test Loss: 0.3807
Epoch 35/50	- Train Loss: 0.1196, Test Loss: 0.3808
Epoch 36/50	- Train Loss: 0.1154, Test Loss: 0.3693
Epoch 37/50	- Train Loss: 0.1076, Test Loss: 0.3690
Epoch 38/50	- Train Loss: 0.1014, Test Loss: 0.3516
Epoch 39/50	- Train Loss: 0.0973, Test Loss: 0.3483
Epoch 40/50	- Train Loss: 0.0912, Test Loss: 0.3546
Epoch 41/50	- Train Loss: 0.0865, Test Loss: 0.3504
Epoch 42/50	- Train Loss: 0.0824, Test Loss: 0.3455
Epoch 43/50	- Train Loss: 0.0800, Test Loss: 0.3346
Epoch 44/50	- Train Loss: 0.0739, Test Loss: 0.3427
Epoch 45/50	- Train Loss: 0.0710, Test Loss: 0.3334
Epoch 46/50	- Train Loss: 0.0676, Test Loss: 0.3292
Epoch 47/50	- Train Loss: 0.0648, Test Loss: 0.3236
Epoch 48/50	- Train Loss: 0.0621, Test Loss: 0.3245
Epoch 49/50	- Train Loss: 0.0587, Test Loss: 0.3160
Epoch 50/50	- Train Loss: 0.0559, Test Loss: 0.3228
Epoch 1/50	- Train Loss: 1.6778, Test Loss: 1.7759
Epoch 2/50	- Train Loss: 1.2531, Test Loss: 1.5112
Epoch 3/50	- Train Loss: 0.9335, Test Loss: 1.2444
Epoch 4/50	- Train Loss: 0.7094, Test Loss: 0.9866
Epoch 5/50	- Train Loss: 0.5439, Test Loss: 0.7196
Epoch 6/50	- Train Loss: 0.4163, Test Loss: 0.8180
Epoch 7/50	- Train Loss: 0.2978, Test Loss: 0.5042
Epoch 8/50	- Train Loss: 0.2246, Test Loss: 0.4446
Epoch 9/50	- Train Loss: 0.1669, Test Loss: 0.5622
Epoch 10/50	- Train Loss: 0.1306, Test Loss: 0.4050
Epoch 11/50	- Train Loss: 0.1020, Test Loss: 0.4054
Epoch 12/50	- Train Loss: 0.0892, Test Loss: 0.5016
Epoch 13/50	- Train Loss: 0.0654, Test Loss: 0.3675
Epoch 14/50	- Train Loss: 0.0554, Test Loss: 0.3410
Epoch 15/50	- Train Loss: 0.0445, Test Loss: 0.3679
Epoch 16/50	- Train Loss: 0.0369, Test Loss: 0.3798
Epoch 17/50	- Train Loss: 0.0302, Test Loss: 0.3623
Epoch 18/50	- Train Loss: 0.0262, Test Loss: 0.3146
Epoch 19/50	- Train Loss: 0.0234, Test Loss: 0.3022
Epoch 20/50	- Train Loss: 0.0194, Test Loss: 0.3095

Epoch 21/50 - Train Loss: 0.0171, Test Loss: 0.3281  
Epoch 22/50 - Train Loss: 0.0148, Test Loss: 0.3165  
Epoch 23/50 - Train Loss: 0.0132, Test Loss: 0.3279  
Epoch 24/50 - Train Loss: 0.0113, Test Loss: 0.3183  
Epoch 25/50 - Train Loss: 0.0105, Test Loss: 0.3264  
Epoch 26/50 - Train Loss: 0.0094, Test Loss: 0.3375  
Epoch 27/50 - Train Loss: 0.0085, Test Loss: 0.3413  
Epoch 28/50 - Train Loss: 0.0078, Test Loss: 0.3371  
Epoch 29/50 - Train Loss: 0.0072, Test Loss: 0.3701  
Epoch 30/50 - Train Loss: 0.0067, Test Loss: 0.3199  
Epoch 31/50 - Train Loss: 0.0061, Test Loss: 0.3304  
Epoch 32/50 - Train Loss: 0.0056, Test Loss: 0.3228  
Epoch 33/50 - Train Loss: 0.0052, Test Loss: 0.3322  
Epoch 34/50 - Train Loss: 0.0049, Test Loss: 0.3309  
Epoch 35/50 - Train Loss: 0.0045, Test Loss: 0.3451  
Epoch 36/50 - Train Loss: 0.0042, Test Loss: 0.3385  
Epoch 37/50 - Train Loss: 0.0039, Test Loss: 0.3665  
Epoch 38/50 - Train Loss: 0.0040, Test Loss: 0.3534  
Epoch 39/50 - Train Loss: 0.0036, Test Loss: 0.3676  
Epoch 40/50 - Train Loss: 0.0034, Test Loss: 0.3364  
Epoch 41/50 - Train Loss: 0.0031, Test Loss: 0.3412  
Epoch 42/50 - Train Loss: 0.0031, Test Loss: 0.3612  
Epoch 43/50 - Train Loss: 0.0028, Test Loss: 0.3486  
Epoch 44/50 - Train Loss: 0.0027, Test Loss: 0.3427  
Epoch 45/50 - Train Loss: 0.0026, Test Loss: 0.3474  
Epoch 46/50 - Train Loss: 0.0024, Test Loss: 0.3392  
Epoch 47/50 - Train Loss: 0.0023, Test Loss: 0.3619  
Epoch 48/50 - Train Loss: 0.0022, Test Loss: 0.3447  
Epoch 49/50 - Train Loss: 0.0021, Test Loss: 0.3608  
Epoch 50/50 - Train Loss: 0.0020, Test Loss: 0.3562  
Epoch 1/50 - Train Loss: 22.8699, Test Loss: 2.3977  
Epoch 2/50 - Train Loss: 1.7445, Test Loss: 1.7409  
Epoch 3/50 - Train Loss: 1.5566, Test Loss: 1.4337  
Epoch 4/50 - Train Loss: 1.3847, Test Loss: 1.2602  
Epoch 5/50 - Train Loss: 1.2266, Test Loss: 1.1817  
Epoch 6/50 - Train Loss: 1.1633, Test Loss: 1.1820  
Epoch 7/50 - Train Loss: 1.1160, Test Loss: 1.1936  
Epoch 8/50 - Train Loss: 1.0604, Test Loss: 1.1813  
Epoch 9/50 - Train Loss: 0.9786, Test Loss: 1.1458  
Epoch 10/50 - Train Loss: 0.8941, Test Loss: 1.1762  
Epoch 11/50 - Train Loss: 0.8522, Test Loss: 1.1564  
Epoch 12/50 - Train Loss: 0.7800, Test Loss: 1.1923  
Epoch 13/50 - Train Loss: 0.7536, Test Loss: 1.1700  
Epoch 14/50 - Train Loss: 0.6717, Test Loss: 1.2439  
Epoch 15/50 - Train Loss: 0.6508, Test Loss: 1.0891  
Epoch 16/50 - Train Loss: 0.6444, Test Loss: 1.1217  
Epoch 17/50 - Train Loss: 0.6297, Test Loss: 1.3271  
Epoch 18/50 - Train Loss: 0.5554, Test Loss: 1.1415  
Epoch 19/50 - Train Loss: 0.6024, Test Loss: 1.3680  
Epoch 20/50 - Train Loss: 0.5529, Test Loss: 1.1485  
Epoch 21/50 - Train Loss: 0.5047, Test Loss: 0.8801  
Epoch 22/50 - Train Loss: 0.5346, Test Loss: 1.1646  
Epoch 23/50 - Train Loss: 0.6528, Test Loss: 0.9063  
Epoch 24/50 - Train Loss: 0.6743, Test Loss: 2.0402  
Epoch 25/50 - Train Loss: 0.5574, Test Loss: 1.2598  
Epoch 26/50 - Train Loss: 0.4491, Test Loss: 1.5941  
Epoch 27/50 - Train Loss: 0.5834, Test Loss: 1.4442  
Epoch 28/50 - Train Loss: 0.7787, Test Loss: 1.7228  
Epoch 29/50 - Train Loss: 0.5801, Test Loss: 1.1907  
Epoch 30/50 - Train Loss: 0.5644, Test Loss: 1.1553

```

Epoch 31/50 - Train Loss: 0.4215, Test Loss: 1.2258
Epoch 32/50 - Train Loss: 0.3945, Test Loss: 1.7066
Epoch 33/50 - Train Loss: 0.3358, Test Loss: 1.4442
Epoch 34/50 - Train Loss: 0.3488, Test Loss: 1.6962
Epoch 35/50 - Train Loss: 0.4235, Test Loss: 1.7088
Epoch 36/50 - Train Loss: 0.3402, Test Loss: 1.3661
Epoch 37/50 - Train Loss: 0.5351, Test Loss: 1.3780
Epoch 38/50 - Train Loss: 0.4855, Test Loss: 1.4712
Epoch 39/50 - Train Loss: 0.5154, Test Loss: 1.6123
Epoch 40/50 - Train Loss: 0.3894, Test Loss: 1.6726
Epoch 41/50 - Train Loss: 0.4135, Test Loss: 1.7034
Epoch 42/50 - Train Loss: 0.4078, Test Loss: 1.3964
Epoch 43/50 - Train Loss: 0.4509, Test Loss: 1.9534
Epoch 44/50 - Train Loss: 0.4024, Test Loss: 2.1903
Epoch 45/50 - Train Loss: 0.3604, Test Loss: 2.0453
Epoch 46/50 - Train Loss: 0.3732, Test Loss: 1.6542
Epoch 47/50 - Train Loss: 0.2737, Test Loss: 1.7814
Epoch 48/50 - Train Loss: 0.2783, Test Loss: 2.1356
Epoch 49/50 - Train Loss: 0.2779, Test Loss: 2.7634
Epoch 50/50 - Train Loss: 0.2830, Test Loss: 2.2819


```

```

In [14]: # TODO: Visualize the results
# Training loss and accuracy
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.title("Training Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
for i, lr in enumerate(learning_rates):
    plt.plot(train_losses[i], label=f"LR: {lr}")
plt.legend()
plt.subplot(1, 2, 2)
plt.title("Training Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
for i, lr in enumerate(learning_rates):
    plt.plot(train_accuracies[i], label=f"LR: {lr}")
plt.legend()
plt.tight_layout()
plt.show()

# Validation loss and accuracy
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.title("Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
for i, lr in enumerate(learning_rates):
    plt.plot(val_losses[i], label=f"LR: {lr}")
plt.legend()
plt.subplot(1, 2, 2)
plt.title("Validation Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
for i, lr in enumerate(learning_rates):
    plt.plot(val_accuracies[i], label=f"LR: {lr}")
plt.legend()
plt.tight_layout()
plt.show()

```



For both extremes in the learning rate (0.1 and 1e-05), one can see, that in the first case the step size is too large to find good parameters in the parameter space. So the training is spiky. For the small learning rate we do not have enough epochs to reach a convergence. The two Learning rates in the middle both converge and reach similar results, while 0.001 converges faster.

## Task 4

```
In [15]: # TODO: apply the best model to the test set
# Take the model with LR = 0.001 as best model because it converges the f

best_model = models[2]

best_model.eval()
test_loss = 0
test_acc = 0
with tc.no_grad():
    for X, Y in test_loader:
        X, Y = X.to(device), Y.to(device)
        output = best_model(X)
        loss = nn.NLLLoss()(output, Y)
        test_loss += loss.item() / len(test_loader)
```

```
test_acc += (output.argmax(dim=1) == Y).float().mean().item() / 1
print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {test_acc:.4f}")
```

Test Loss: 0.1991, Test Accuracy: 0.9250

The Validation and Test error have around the same accuracy, while the loss is much smaller in the validation loss.

## Exercise 3 - CNN Autoencoder

In the next task we want to build an autoencoder. It consists of an encoder, which transforms the data into a low-dimensional code, and a decoder, which reconstructs the original data.

We use the Fashion MNIST dataset, which consists of 28x28 grayscale images. There are 10 classes, each representing different items of clothing. The data can be conveniently downloaded, separated and transformed with torchvision. As we will not be tuning any hyperparameters, we do not need a validation set.

```
In [96]: from torchvision import datasets, transforms

# Define transformations
transform = transforms.Compose([
    transforms.ToTensor(), # Convert PIL image to tensor (range [0, 1])
    transforms.Normalize((0.5,), (0.5,)) # Normalize to range [-1, 1]
])

# Load FashionMNIST train and test sets
train_data = datasets.FashionMNIST(
    root='./fashion_mnist', # Download path
    train=True, # Load training set
    download=True, # Download if not already present
    transform=transform # Apply transformations
)

test_data = datasets.FashionMNIST(
    root='./fashion_mnist',
    train=False, # Load test set
    download=True,
    transform=transform
)

classes = train_data.targets.unique()

# we only need a subset that consists of 1000 samples of each class for t
# and 10 samples of each class for the test set
indices_train = []
indices_test = []
for i in range(len(classes)):
    indices_train += list(np.where(train_data.targets == classes[i])[0][:
    indices_test += list(np.where(test_data.targets == classes[i])[0][:10

train_data.data = train_data.data[indices_train]
train_data.targets = train_data.targets[indices_train]
test_data.data = test_data.data[indices_test]
test_data.targets = test_data.targets[indices_test]
```

```
# Create DataLoaders
train_batch_size = 256
test_batch_size = len(test_data)
train_loader = DataLoader(train_data, batch_size=train_batch_size, shuffle=True)
test_loader = DataLoader(test_data, batch_size=test_batch_size, shuffle=False)


# Print the shape of the dataset
print("train images shape: " + str(train_data.data.shape))
print("train labels shape: " + str(train_data.targets.shape))
print("test images shape: " + str(test_data.data.shape))
print("test labels shape: " + str(test_data.targets.shape))
print("classes: " + str(classes))

# check if classes are balanced
print("Counts of classes in train set: " + str(train_data.targets.unique(return_counts=True)[1]))
print("Counts of classes in test set: " + str(test_data.targets.unique(return_counts=True)[1]))
```

```
train images shape: torch.Size([10000, 28, 28])
train labels shape: torch.Size([10000])
test images shape: torch.Size([100, 28, 28])
test labels shape: torch.Size([100])
classes: tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
Counts of classes in train set: tensor([1000, 1000, 1000, 1000, 1000, 1000, 1000, 1000, 1000, 1000])
Counts of classes in test set: tensor([10, 10, 10, 10, 10, 10, 10, 10, 10, 10])
```

```
In [30]: # Get a batch of data
data_iter = iter(train_loader)
images, labels = next(data_iter)

# Visualize a batch of images
fig, axes = plt.subplots(4, 10, figsize=(12, 6))
for i in range(4):
    for j in range(10):
        # find the ith image of class j
        idx = np.where(labels == j)[0][i]
        axes[i, j].imshow(images[idx].squeeze(), cmap='gray')
        axes[i, j].set_title(f"Label: {j}")
        axes[i, j].axis('off')
plt.tight_layout()
plt.show()
```



## Task 1

We compare two architectures: a linear autoencoder and a CNN autoencoder. The latter typically consists of convolutional layers for the encoder and transposed convolutional layers for the decoder. In addition, fully connected layers can bring the feature to the desired code dimension (also called latent dimension). Implement the following architecture:

Encoder:

- Conv2d: 16 output channels, 3 by 3 filter size, stride 1, padding "same"
- ReLU activation
- MaxPool2d: 2 by 2 filter, stride 1, padding 0
- Conv2d: 32 output channels, 3 by 3, stride 1, padding "same"
- ReLU activation
- MaxPool2d: 2 by 2 filter, stride 1, padding 0
- Conv2d: 64 output channel, 3 by 3 filter size, stride 1, padding "same"
- ReLU activation
- MaxPool2d: 2 by 2 filter, stride 1, padding 0
- Flatten the previous output



- Linear: *<latent dimension>* output neurons

Decoder:

- Linear:  $64 \times 3 \times 3 = 576$  output neurons
- Unflatten the previous output to shape (64, 3, 3)
- ConvTranspose2d: 32 output channels, 3 by 3 filter size, stride 2, padding 0, output padding 1
- ReLU activation
- ConvTranspose2d: 16 output channels, 3 by 3 filter size, stride 2, padding 1, output padding 1
- ReLU activation
- ConvTranspose2d: 1 output channel, 3 by 3 filter size, stride 2, padding 1, output padding 1

You might need to infer the input dimension of the linear layer in the encoder.

```
In [59]: class Conv_AE(nn.Module):
    def __init__(self, latent_dim):
        super(Conv_AE, self).__init__()

        # TODO: Initialize the layers of the encoder
        self.encoder = nn.Sequential(
            nn.Conv2d(1, 16, (3,3), 1, 'same'),
            nn.ReLU(),
            nn.MaxPool2d((2,2), 1),
            # output dim= 27
            nn.Conv2d(16, 32, (3,3), 1, 'same'),
            nn.ReLU(),
            nn.MaxPool2d((2,2), 1, 0),
            # output dim = 26
            nn.Conv2d(32, 64, (3,3), 1, 'same'),
            nn.ReLU(),
            nn.MaxPool2d((2,2), 1, 0),
            # output dim = 25
            nn.Flatten(),
            # output dim = 25*25*64
            nn.Linear(40000, latent_dim)
        )

        # TODO: Initialize the layers of the decoder
        self.decoder = nn.Sequential(
            nn.Linear(latent_dim, 576),
            nn.Unflatten(1, (64, 3, 3)),
            nn.ConvTranspose2d(64, 32, (3,3), 2, 0, 0),
            nn.ReLU(),
            nn.ConvTranspose2d(32, 16, (3,3), 2, 1, 1),
            nn.ReLU(),
            nn.ConvTranspose2d(16, 1, (3,3), 2, 1, 1)
        )

    def forward(self, X):
        # TODO: Implement the forward pass
        latent = self.encoder(X)

        return self.decoder(latent), latent
```

We check again whether the model is working properly.

```
In [60]: conv_ae_model = Conv_AE(2)
X_random = tc.randn(2, 1, 28, 28)
reconstructed, latent = conv_ae_model(X_random)

print("Reconstructed shape: " + str(reconstructed.shape), "\n", "Latent s
summary(conv_ae_model, input_size=(1, 28, 28), device="cpu")
```

Reconstructed shape: torch.Size([2, 1, 28, 28])

Latent shape: torch.Size([2, 2])

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 28, 28]	160
ReLU-2	[-1, 16, 28, 28]	0
MaxPool2d-3	[-1, 16, 27, 27]	0
Conv2d-4	[-1, 32, 27, 27]	4,640
ReLU-5	[-1, 32, 27, 27]	0
MaxPool2d-6	[-1, 32, 26, 26]	0
Conv2d-7	[-1, 64, 26, 26]	18,496
ReLU-8	[-1, 64, 26, 26]	0
MaxPool2d-9	[-1, 64, 25, 25]	0
Flatten-10	[-1, 40000]	0
Linear-11	[-1, 2]	80,002
Linear-12	[-1, 576]	1,728
Unflatten-13	[-1, 64, 3, 3]	0
ConvTranspose2d-14	[-1, 32, 7, 7]	18,464
ReLU-15	[-1, 32, 7, 7]	0
ConvTranspose2d-16	[-1, 16, 14, 14]	4,624
ReLU-17	[-1, 16, 14, 14]	0
ConvTranspose2d-18	[-1, 1, 28, 28]	145

Total params: 128,259

Trainable params: 128,259

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 2.16

Params size (MB): 0.49

Estimated Total Size (MB): 2.65

## Task 2

The training of an autoencoder compares the original input to the reconstruction, usually by means of the mean squared error.

```
In [61]: def train_ae(model, train_loader, test_loader, lr, n_epochs, device):
    model = model.to(device)

    # TODO: Initialize the loss function
    loss_function = nn.MSELoss()
    optimizer = optim.Adam(model.parameters(), lr)

    train_loss = np.zeros(n_epochs)
    test_loss = np.zeros(n_epochs)

    for epoch in range(1, n_epochs + 1):
        model.train()

        epoch_loss = 0

        for X, _ in train_loader:
            X = X.to(device)
            # TODO: Implement the training loop
            optimizer.zero_grad()
```

```

        output, _ = model(X)

        loss = loss_function(X, output)

        loss.backward()
        optimizer.step()

        epoch_loss += loss.item()/len(train_loader)

    train_loss[epoch - 1] = epoch_loss

    model.eval()

    with tc.no_grad():
        X, _ = next(iter(test_loader))
        X = X.to(device)
        # TODO: Implement the evaluation step
        output, _ = model(X)
        loss = loss_function(X, output)

        test_loss[epoch - 1] = loss.item()
    print(f"Epoch {epoch}/{n_epochs} - Train Loss: {epoch_loss:.4f},

    return train_loss, test_loss

```

```

In [62]: n_epochs = 60
        lr = 1e-3
        # TODO: Train the convolutional autoencoder model with different latent d
        latent_dims = [3, 10, 50, 100, 250]

        models = [Conv_AE(i) for i in latent_dims]
        train_losses = []
        test_losses = []
        for latent_dim in latent_dims:
            model = models[latent_dims.index(latent_dim)]
            train_loss, test_loss = train_ae(model, train_loader, test_loader, lr
            train_losses.append(train_loss)
            test_losses.append(test_loss)

```

Epoch 1/60 - Train Loss: 0.3350, Test Loss: 0.1932  
Epoch 2/60 - Train Loss: 0.1757, Test Loss: 0.1613  
Epoch 3/60 - Train Loss: 0.1577, Test Loss: 0.1504  
Epoch 4/60 - Train Loss: 0.1460, Test Loss: 0.1429  
Epoch 5/60 - Train Loss: 0.1386, Test Loss: 0.1356  
Epoch 6/60 - Train Loss: 0.1336, Test Loss: 0.1280  
Epoch 7/60 - Train Loss: 0.1284, Test Loss: 0.1234  
Epoch 8/60 - Train Loss: 0.1240, Test Loss: 0.1175  
Epoch 9/60 - Train Loss: 0.1213, Test Loss: 0.1146  
Epoch 10/60 - Train Loss: 0.1186, Test Loss: 0.1134  
Epoch 11/60 - Train Loss: 0.1167, Test Loss: 0.1133  
Epoch 12/60 - Train Loss: 0.1150, Test Loss: 0.1120  
Epoch 13/60 - Train Loss: 0.1144, Test Loss: 0.1090  
Epoch 14/60 - Train Loss: 0.1131, Test Loss: 0.1082  
Epoch 15/60 - Train Loss: 0.1121, Test Loss: 0.1074  
Epoch 16/60 - Train Loss: 0.1108, Test Loss: 0.1060  
Epoch 17/60 - Train Loss: 0.1102, Test Loss: 0.1069  
Epoch 18/60 - Train Loss: 0.1097, Test Loss: 0.1049  
Epoch 19/60 - Train Loss: 0.1097, Test Loss: 0.1054  
Epoch 20/60 - Train Loss: 0.1091, Test Loss: 0.1047  
Epoch 21/60 - Train Loss: 0.1092, Test Loss: 0.1070  
Epoch 22/60 - Train Loss: 0.1086, Test Loss: 0.1034  
Epoch 23/60 - Train Loss: 0.1076, Test Loss: 0.1037  
Epoch 24/60 - Train Loss: 0.1075, Test Loss: 0.1031  
Epoch 25/60 - Train Loss: 0.1064, Test Loss: 0.1026  
Epoch 26/60 - Train Loss: 0.1065, Test Loss: 0.1047  
Epoch 27/60 - Train Loss: 0.1072, Test Loss: 0.1021  
Epoch 28/60 - Train Loss: 0.1052, Test Loss: 0.1027  
Epoch 29/60 - Train Loss: 0.1055, Test Loss: 0.1023  
Epoch 30/60 - Train Loss: 0.1062, Test Loss: 0.1039  
Epoch 31/60 - Train Loss: 0.1063, Test Loss: 0.1010  
Epoch 32/60 - Train Loss: 0.1041, Test Loss: 0.1036  
Epoch 33/60 - Train Loss: 0.1043, Test Loss: 0.1002  
Epoch 34/60 - Train Loss: 0.1043, Test Loss: 0.1025  
Epoch 35/60 - Train Loss: 0.1044, Test Loss: 0.1017  
Epoch 36/60 - Train Loss: 0.1045, Test Loss: 0.1017  
Epoch 37/60 - Train Loss: 0.1035, Test Loss: 0.1007  
Epoch 38/60 - Train Loss: 0.1029, Test Loss: 0.0993  
Epoch 39/60 - Train Loss: 0.1025, Test Loss: 0.0993  
Epoch 40/60 - Train Loss: 0.1027, Test Loss: 0.0994  
Epoch 41/60 - Train Loss: 0.1023, Test Loss: 0.0999  
Epoch 42/60 - Train Loss: 0.1024, Test Loss: 0.1011  
Epoch 43/60 - Train Loss: 0.1022, Test Loss: 0.0996  
Epoch 44/60 - Train Loss: 0.1022, Test Loss: 0.1005  
Epoch 45/60 - Train Loss: 0.1017, Test Loss: 0.1001  
Epoch 46/60 - Train Loss: 0.1015, Test Loss: 0.0983  
Epoch 47/60 - Train Loss: 0.1017, Test Loss: 0.1012  
Epoch 48/60 - Train Loss: 0.1012, Test Loss: 0.0987  
Epoch 49/60 - Train Loss: 0.1000, Test Loss: 0.0983  
Epoch 50/60 - Train Loss: 0.1002, Test Loss: 0.0991  
Epoch 51/60 - Train Loss: 0.1005, Test Loss: 0.0978  
Epoch 52/60 - Train Loss: 0.1004, Test Loss: 0.0977  
Epoch 53/60 - Train Loss: 0.1007, Test Loss: 0.0997  
Epoch 54/60 - Train Loss: 0.1005, Test Loss: 0.0999  
Epoch 55/60 - Train Loss: 0.1005, Test Loss: 0.0985  
Epoch 56/60 - Train Loss: 0.0996, Test Loss: 0.0994  
Epoch 57/60 - Train Loss: 0.0993, Test Loss: 0.0977  
Epoch 58/60 - Train Loss: 0.0991, Test Loss: 0.0972  
Epoch 59/60 - Train Loss: 0.0995, Test Loss: 0.0973  
Epoch 60/60 - Train Loss: 0.0991, Test Loss: 0.0968

Epoch 1/60 - Train Loss: 0.4024, Test Loss: 0.2187  
Epoch 2/60 - Train Loss: 0.1651, Test Loss: 0.1288  
Epoch 3/60 - Train Loss: 0.1202, Test Loss: 0.1101  
Epoch 4/60 - Train Loss: 0.1063, Test Loss: 0.0995  
Epoch 5/60 - Train Loss: 0.0988, Test Loss: 0.0945  
Epoch 6/60 - Train Loss: 0.0939, Test Loss: 0.0894  
Epoch 7/60 - Train Loss: 0.0911, Test Loss: 0.0861  
Epoch 8/60 - Train Loss: 0.0879, Test Loss: 0.0839  
Epoch 9/60 - Train Loss: 0.0857, Test Loss: 0.0821  
Epoch 10/60 - Train Loss: 0.0837, Test Loss: 0.0804  
Epoch 11/60 - Train Loss: 0.0821, Test Loss: 0.0799  
Epoch 12/60 - Train Loss: 0.0812, Test Loss: 0.0779  
Epoch 13/60 - Train Loss: 0.0803, Test Loss: 0.0775  
Epoch 14/60 - Train Loss: 0.0792, Test Loss: 0.0768  
Epoch 15/60 - Train Loss: 0.0781, Test Loss: 0.0752  
Epoch 16/60 - Train Loss: 0.0767, Test Loss: 0.0742  
Epoch 17/60 - Train Loss: 0.0765, Test Loss: 0.0741  
Epoch 18/60 - Train Loss: 0.0757, Test Loss: 0.0736  
Epoch 19/60 - Train Loss: 0.0749, Test Loss: 0.0731  
Epoch 20/60 - Train Loss: 0.0742, Test Loss: 0.0720  
Epoch 21/60 - Train Loss: 0.0736, Test Loss: 0.0719  
Epoch 22/60 - Train Loss: 0.0730, Test Loss: 0.0707  
Epoch 23/60 - Train Loss: 0.0727, Test Loss: 0.0711  
Epoch 24/60 - Train Loss: 0.0727, Test Loss: 0.0701  
Epoch 25/60 - Train Loss: 0.0718, Test Loss: 0.0702  
Epoch 26/60 - Train Loss: 0.0715, Test Loss: 0.0696  
Epoch 27/60 - Train Loss: 0.0706, Test Loss: 0.0687  
Epoch 28/60 - Train Loss: 0.0704, Test Loss: 0.0682  
Epoch 29/60 - Train Loss: 0.0704, Test Loss: 0.0690  
Epoch 30/60 - Train Loss: 0.0698, Test Loss: 0.0678  
Epoch 31/60 - Train Loss: 0.0691, Test Loss: 0.0681  
Epoch 32/60 - Train Loss: 0.0693, Test Loss: 0.0675  
Epoch 33/60 - Train Loss: 0.0690, Test Loss: 0.0669  
Epoch 34/60 - Train Loss: 0.0681, Test Loss: 0.0668  
Epoch 35/60 - Train Loss: 0.0676, Test Loss: 0.0664  
Epoch 36/60 - Train Loss: 0.0681, Test Loss: 0.0665  
Epoch 37/60 - Train Loss: 0.0675, Test Loss: 0.0664  
Epoch 38/60 - Train Loss: 0.0681, Test Loss: 0.0663  
Epoch 39/60 - Train Loss: 0.0673, Test Loss: 0.0658  
Epoch 40/60 - Train Loss: 0.0668, Test Loss: 0.0657  
Epoch 41/60 - Train Loss: 0.0665, Test Loss: 0.0650  
Epoch 42/60 - Train Loss: 0.0659, Test Loss: 0.0649  
Epoch 43/60 - Train Loss: 0.0658, Test Loss: 0.0646  
Epoch 44/60 - Train Loss: 0.0660, Test Loss: 0.0653  
Epoch 45/60 - Train Loss: 0.0656, Test Loss: 0.0649  
Epoch 46/60 - Train Loss: 0.0653, Test Loss: 0.0644  
Epoch 47/60 - Train Loss: 0.0652, Test Loss: 0.0644  
Epoch 48/60 - Train Loss: 0.0654, Test Loss: 0.0643  
Epoch 49/60 - Train Loss: 0.0654, Test Loss: 0.0641  
Epoch 50/60 - Train Loss: 0.0648, Test Loss: 0.0638  
Epoch 51/60 - Train Loss: 0.0648, Test Loss: 0.0638  
Epoch 52/60 - Train Loss: 0.0644, Test Loss: 0.0638  
Epoch 53/60 - Train Loss: 0.0644, Test Loss: 0.0633  
Epoch 54/60 - Train Loss: 0.0641, Test Loss: 0.0633  
Epoch 55/60 - Train Loss: 0.0643, Test Loss: 0.0632  
Epoch 56/60 - Train Loss: 0.0641, Test Loss: 0.0630  
Epoch 57/60 - Train Loss: 0.0633, Test Loss: 0.0628  
Epoch 58/60 - Train Loss: 0.0631, Test Loss: 0.0626  
Epoch 59/60 - Train Loss: 0.0632, Test Loss: 0.0625  
Epoch 60/60 - Train Loss: 0.0630, Test Loss: 0.0631


```
Epoch 1/60 - Train Loss: 0.3719, Test Loss: 0.2227
Epoch 2/60 - Train Loss: 0.1554, Test Loss: 0.1114
Epoch 3/60 - Train Loss: 0.0994, Test Loss: 0.0903
Epoch 4/60 - Train Loss: 0.0834, Test Loss: 0.0772
Epoch 5/60 - Train Loss: 0.0745, Test Loss: 0.0696
Epoch 6/60 - Train Loss: 0.0689, Test Loss: 0.0651
Epoch 7/60 - Train Loss: 0.0647, Test Loss: 0.0631
Epoch 8/60 - Train Loss: 0.0614, Test Loss: 0.0590
Epoch 9/60 - Train Loss: 0.0585, Test Loss: 0.0562
Epoch 10/60 - Train Loss: 0.0563, Test Loss: 0.0555
Epoch 11/60 - Train Loss: 0.0540, Test Loss: 0.0546
Epoch 12/60 - Train Loss: 0.0536, Test Loss: 0.0514
Epoch 13/60 - Train Loss: 0.0513, Test Loss: 0.0500
Epoch 14/60 - Train Loss: 0.0499, Test Loss: 0.0489
Epoch 15/60 - Train Loss: 0.0490, Test Loss: 0.0482
Epoch 16/60 - Train Loss: 0.0479, Test Loss: 0.0467
Epoch 17/60 - Train Loss: 0.0467, Test Loss: 0.0458
Epoch 18/60 - Train Loss: 0.0461, Test Loss: 0.0452
Epoch 19/60 - Train Loss: 0.0452, Test Loss: 0.0441
Epoch 20/60 - Train Loss: 0.0440, Test Loss: 0.0434
Epoch 21/60 - Train Loss: 0.0438, Test Loss: 0.0434
Epoch 22/60 - Train Loss: 0.0432, Test Loss: 0.0420
Epoch 23/60 - Train Loss: 0.0420, Test Loss: 0.0418
Epoch 24/60 - Train Loss: 0.0415, Test Loss: 0.0412
Epoch 25/60 - Train Loss: 0.0410, Test Loss: 0.0403
Epoch 26/60 - Train Loss: 0.0407, Test Loss: 0.0396
Epoch 27/60 - Train Loss: 0.0398, Test Loss: 0.0389
Epoch 28/60 - Train Loss: 0.0394, Test Loss: 0.0388
Epoch 29/60 - Train Loss: 0.0390, Test Loss: 0.0384
Epoch 30/60 - Train Loss: 0.0385, Test Loss: 0.0388
Epoch 31/60 - Train Loss: 0.0384, Test Loss: 0.0374
Epoch 32/60 - Train Loss: 0.0374, Test Loss: 0.0369
Epoch 33/60 - Train Loss: 0.0369, Test Loss: 0.0366
Epoch 34/60 - Train Loss: 0.0367, Test Loss: 0.0367
Epoch 35/60 - Train Loss: 0.0363, Test Loss: 0.0359
Epoch 36/60 - Train Loss: 0.0358, Test Loss: 0.0358
Epoch 37/60 - Train Loss: 0.0358, Test Loss: 0.0359
Epoch 38/60 - Train Loss: 0.0354, Test Loss: 0.0358
Epoch 39/60 - Train Loss: 0.0349, Test Loss: 0.0350
Epoch 40/60 - Train Loss: 0.0347, Test Loss: 0.0350
Epoch 41/60 - Train Loss: 0.0345, Test Loss: 0.0347
Epoch 42/60 - Train Loss: 0.0342, Test Loss: 0.0340
Epoch 43/60 - Train Loss: 0.0342, Test Loss: 0.0347
Epoch 44/60 - Train Loss: 0.0338, Test Loss: 0.0336
Epoch 45/60 - Train Loss: 0.0333, Test Loss: 0.0342
Epoch 46/60 - Train Loss: 0.0333, Test Loss: 0.0338
Epoch 47/60 - Train Loss: 0.0329, Test Loss: 0.0336
Epoch 48/60 - Train Loss: 0.0325, Test Loss: 0.0329
Epoch 49/60 - Train Loss: 0.0325, Test Loss: 0.0332
Epoch 50/60 - Train Loss: 0.0323, Test Loss: 0.0327
Epoch 51/60 - Train Loss: 0.0320, Test Loss: 0.0331
Epoch 52/60 - Train Loss: 0.0321, Test Loss: 0.0333
Epoch 53/60 - Train Loss: 0.0317, Test Loss: 0.0325
Epoch 54/60 - Train Loss: 0.0320, Test Loss: 0.0328
Epoch 55/60 - Train Loss: 0.0317, Test Loss: 0.0326
Epoch 56/60 - Train Loss: 0.0314, Test Loss: 0.0324
Epoch 57/60 - Train Loss: 0.0309, Test Loss: 0.0321
Epoch 58/60 - Train Loss: 0.0309, Test Loss: 0.0320
Epoch 59/60 - Train Loss: 0.0307, Test Loss: 0.0320
Epoch 60/60 - Train Loss: 0.0305, Test Loss: 0.0320
```

Epoch 1/60 - Train Loss: 0.3333, Test Loss: 0.1486  
Epoch 2/60 - Train Loss: 0.1203, Test Loss: 0.0985  
Epoch 3/60 - Train Loss: 0.0900, Test Loss: 0.0802  
Epoch 4/60 - Train Loss: 0.0766, Test Loss: 0.0708  
Epoch 5/60 - Train Loss: 0.0686, Test Loss: 0.0635  
Epoch 6/60 - Train Loss: 0.0629, Test Loss: 0.0605  
Epoch 7/60 - Train Loss: 0.0583, Test Loss: 0.0554  
Epoch 8/60 - Train Loss: 0.0547, Test Loss: 0.0517  
Epoch 9/60 - Train Loss: 0.0552, Test Loss: 0.0506  
Epoch 10/60 - Train Loss: 0.0495, Test Loss: 0.0473  
Epoch 11/60 - Train Loss: 0.0474, Test Loss: 0.0455  
Epoch 12/60 - Train Loss: 0.0457, Test Loss: 0.0443  
Epoch 13/60 - Train Loss: 0.0442, Test Loss: 0.0427  
Epoch 14/60 - Train Loss: 0.0423, Test Loss: 0.0421  
Epoch 15/60 - Train Loss: 0.0415, Test Loss: 0.0399  
Epoch 16/60 - Train Loss: 0.0398, Test Loss: 0.0395  
Epoch 17/60 - Train Loss: 0.0391, Test Loss: 0.0378  
Epoch 18/60 - Train Loss: 0.0378, Test Loss: 0.0371  
Epoch 19/60 - Train Loss: 0.0370, Test Loss: 0.0369  
Epoch 20/60 - Train Loss: 0.0366, Test Loss: 0.0358  
Epoch 21/60 - Train Loss: 0.0354, Test Loss: 0.0347  
Epoch 22/60 - Train Loss: 0.0345, Test Loss: 0.0340  
Epoch 23/60 - Train Loss: 0.0337, Test Loss: 0.0331  
Epoch 24/60 - Train Loss: 0.0332, Test Loss: 0.0324  
Epoch 25/60 - Train Loss: 0.0328, Test Loss: 0.0320  
Epoch 26/60 - Train Loss: 0.0323, Test Loss: 0.0315  
Epoch 27/60 - Train Loss: 0.0313, Test Loss: 0.0311  
Epoch 28/60 - Train Loss: 0.0312, Test Loss: 0.0310  
Epoch 29/60 - Train Loss: 0.0305, Test Loss: 0.0302  
Epoch 30/60 - Train Loss: 0.0301, Test Loss: 0.0296  
Epoch 31/60 - Train Loss: 0.0296, Test Loss: 0.0293  
Epoch 32/60 - Train Loss: 0.0292, Test Loss: 0.0296  
Epoch 33/60 - Train Loss: 0.0294, Test Loss: 0.0285  
Epoch 34/60 - Train Loss: 0.0283, Test Loss: 0.0281  
Epoch 35/60 - Train Loss: 0.0279, Test Loss: 0.0279  
Epoch 36/60 - Train Loss: 0.0279, Test Loss: 0.0275  
Epoch 37/60 - Train Loss: 0.0274, Test Loss: 0.0267  
Epoch 38/60 - Train Loss: 0.0271, Test Loss: 0.0271  
Epoch 39/60 - Train Loss: 0.0265, Test Loss: 0.0266  
Epoch 40/60 - Train Loss: 0.0266, Test Loss: 0.0276  
Epoch 41/60 - Train Loss: 0.0264, Test Loss: 0.0260  
Epoch 42/60 - Train Loss: 0.0255, Test Loss: 0.0257  
Epoch 43/60 - Train Loss: 0.0254, Test Loss: 0.0257  
Epoch 44/60 - Train Loss: 0.0251, Test Loss: 0.0253  
Epoch 45/60 - Train Loss: 0.0248, Test Loss: 0.0251  
Epoch 46/60 - Train Loss: 0.0247, Test Loss: 0.0248  
Epoch 47/60 - Train Loss: 0.0244, Test Loss: 0.0249  
Epoch 48/60 - Train Loss: 0.0243, Test Loss: 0.0246  
Epoch 49/60 - Train Loss: 0.0240, Test Loss: 0.0245  
Epoch 50/60 - Train Loss: 0.0238, Test Loss: 0.0245  
Epoch 51/60 - Train Loss: 0.0237, Test Loss: 0.0240  
Epoch 52/60 - Train Loss: 0.0234, Test Loss: 0.0243  
Epoch 53/60 - Train Loss: 0.0232, Test Loss: 0.0240  
Epoch 54/60 - Train Loss: 0.0230, Test Loss: 0.0241  
Epoch 55/60 - Train Loss: 0.0229, Test Loss: 0.0238  
Epoch 56/60 - Train Loss: 0.0229, Test Loss: 0.0238  
Epoch 57/60 - Train Loss: 0.0227, Test Loss: 0.0236  
Epoch 58/60 - Train Loss: 0.0227, Test Loss: 0.0235  
Epoch 59/60 - Train Loss: 0.0222, Test Loss: 0.0233  
Epoch 60/60 - Train Loss: 0.0223, Test Loss: 0.0232



Epoch 1/60 - Train Loss: 0.3049, Test Loss: 0.1403  
Epoch 2/60 - Train Loss: 0.1108, Test Loss: 0.0917  
Epoch 3/60 - Train Loss: 0.0819, Test Loss: 0.0735  
Epoch 4/60 - Train Loss: 0.0693, Test Loss: 0.0662  
Epoch 5/60 - Train Loss: 0.0614, Test Loss: 0.0567  
Epoch 6/60 - Train Loss: 0.0549, Test Loss: 0.0526  
Epoch 7/60 - Train Loss: 0.0503, Test Loss: 0.0483  
Epoch 8/60 - Train Loss: 0.0471, Test Loss: 0.0450  
Epoch 9/60 - Train Loss: 0.0448, Test Loss: 0.0437  
Epoch 10/60 - Train Loss: 0.0419, Test Loss: 0.0405  
Epoch 11/60 - Train Loss: 0.0398, Test Loss: 0.0383  
Epoch 12/60 - Train Loss: 0.0378, Test Loss: 0.0373  
Epoch 13/60 - Train Loss: 0.0366, Test Loss: 0.0351  
Epoch 14/60 - Train Loss: 0.0350, Test Loss: 0.0336  
Epoch 15/60 - Train Loss: 0.0338, Test Loss: 0.0336  
Epoch 16/60 - Train Loss: 0.0328, Test Loss: 0.0316  
Epoch 17/60 - Train Loss: 0.0313, Test Loss: 0.0308  
Epoch 18/60 - Train Loss: 0.0304, Test Loss: 0.0303  
Epoch 19/60 - Train Loss: 0.0295, Test Loss: 0.0304  
Epoch 20/60 - Train Loss: 0.0289, Test Loss: 0.0279  
Epoch 21/60 - Train Loss: 0.0277, Test Loss: 0.0271  
Epoch 22/60 - Train Loss: 0.0269, Test Loss: 0.0267  
Epoch 23/60 - Train Loss: 0.0265, Test Loss: 0.0279  
Epoch 24/60 - Train Loss: 0.0264, Test Loss: 0.0256  
Epoch 25/60 - Train Loss: 0.0251, Test Loss: 0.0248  
Epoch 26/60 - Train Loss: 0.0246, Test Loss: 0.0247  
Epoch 27/60 - Train Loss: 0.0241, Test Loss: 0.0242  
Epoch 28/60 - Train Loss: 0.0234, Test Loss: 0.0240  
Epoch 29/60 - Train Loss: 0.0232, Test Loss: 0.0263  
Epoch 30/60 - Train Loss: 0.0231, Test Loss: 0.0229  
Epoch 31/60 - Train Loss: 0.0222, Test Loss: 0.0227  
Epoch 32/60 - Train Loss: 0.0217, Test Loss: 0.0222  
Epoch 33/60 - Train Loss: 0.0213, Test Loss: 0.0219  
Epoch 34/60 - Train Loss: 0.0215, Test Loss: 0.0221  
Epoch 35/60 - Train Loss: 0.0208, Test Loss: 0.0215  
Epoch 36/60 - Train Loss: 0.0203, Test Loss: 0.0211  
Epoch 37/60 - Train Loss: 0.0202, Test Loss: 0.0209  
Epoch 38/60 - Train Loss: 0.0199, Test Loss: 0.0222  
Epoch 39/60 - Train Loss: 0.0200, Test Loss: 0.0204  
Epoch 40/60 - Train Loss: 0.0190, Test Loss: 0.0200  
Epoch 41/60 - Train Loss: 0.0189, Test Loss: 0.0204  
Epoch 42/60 - Train Loss: 0.0187, Test Loss: 0.0196  
Epoch 43/60 - Train Loss: 0.0182, Test Loss: 0.0195  
Epoch 44/60 - Train Loss: 0.0182, Test Loss: 0.0191  
Epoch 45/60 - Train Loss: 0.0180, Test Loss: 0.0193  
Epoch 46/60 - Train Loss: 0.0176, Test Loss: 0.0188  
Epoch 47/60 - Train Loss: 0.0173, Test Loss: 0.0190  
Epoch 48/60 - Train Loss: 0.0172, Test Loss: 0.0185  
Epoch 49/60 - Train Loss: 0.0171, Test Loss: 0.0192  
Epoch 50/60 - Train Loss: 0.0168, Test Loss: 0.0183  
Epoch 51/60 - Train Loss: 0.0166, Test Loss: 0.0186  
Epoch 52/60 - Train Loss: 0.0164, Test Loss: 0.0180  
Epoch 53/60 - Train Loss: 0.0160, Test Loss: 0.0177  
Epoch 54/60 - Train Loss: 0.0160, Test Loss: 0.0186  
Epoch 55/60 - Train Loss: 0.0161, Test Loss: 0.0180  
Epoch 56/60 - Train Loss: 0.0156, Test Loss: 0.0175  
Epoch 57/60 - Train Loss: 0.0152, Test Loss: 0.0174  
Epoch 58/60 - Train Loss: 0.0153, Test Loss: 0.0175  
Epoch 59/60 - Train Loss: 0.0150, Test Loss: 0.0174  
Epoch 60/60 - Train Loss: 0.0148, Test Loss: 0.0173

```
In [52]: # TODO: Visualize the results
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.title("Training Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
for i, latent_dim in enumerate(latent_dims):
    plt.plot(train_losses[i], label=f"Latent Dim: {latent_dim}")
plt.legend()
plt.subplot(1, 2, 2)
plt.title("Test Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
for i, latent_dim in enumerate(latent_dims):
    plt.plot(test_losses[i], label=f"Latent Dim: {latent_dim}")
plt.legend()
plt.tight_layout()
plt.show()
```



```
In [69]: # TODO: Visualize the reconstructed images and the original images
def visualize_reconstruction(model, data_loader, device):
    model.eval()
    with tc.no_grad():
        X, _ = next(iter(data_loader))
        X = X.to(device)
        reconstructed, _ = model(X)

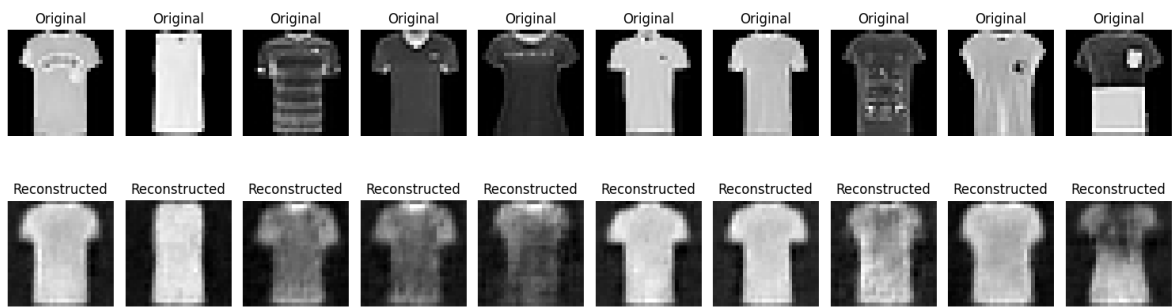
    fig, axs = plt.subplots(2, 10, figsize=(15, 5))
    for i in range(10):
        axs[0, i].imshow(X[i].cpu().squeeze(), cmap='gray')
        axs[0, i].set_title("Original")
        axs[0, i].axis('off')

        axs[1, i].imshow(reconstructed[i].cpu().squeeze(), cmap='gray')
        axs[1, i].set_title("Reconstructed")
        axs[1, i].axis('off')

    plt.tight_layout()
    plt.show()

# Visualize the reconstruction for the model with latent dimension 10
```

```
visualize_reconstruction(models[1], test_loader, device)
```



The shapes of the different fashion pieces is visible but the reconstructed image loses a lot of details.

### Task 3

A linear autoencoder aims to represent the data  $X \in \mathbb{R}^{n \times m}$  in a new basis using only  $d < m$  directions. The objective is to minimize the squared error between  $X$  and  $D(E(X))$  where  $E : \mathbb{R}^m \rightarrow \mathbb{R}^d$  is the encoder and  $D : \mathbb{R}^d \rightarrow \mathbb{R}^m$  is the decoder:

$$\|D(E(X)) - X\|_2^2$$

In geometric terms, we want to find  $d$  axes along which most of the variance occurs which is exactly what Principal Component Analysis does. The optimal weights of a linear autoencoder with code dimension  $d$  thus span the same space as the first  $d$  principal components.

The PCA autoencoder just consists of 1 linear layer for the encoder and 1 linear layer for the decoder. The bias is necessary to subtract the mean value. Since we are dealing with three-dimensional images, we also have to flatten (when encoding) or unflatten (when decoding) the input.

```
In [65]: class PCA_AE(nn.Module):
    def __init__(self, input_dim, latent_dim):
        super(PCA_AE, self).__init__()

        # TODO: Initialize the layers of the autoencoder
        self.encoder = nn.Sequential(
            nn.Flatten(),
            nn.Linear(input_dim, latent_dim)
        )

        self.decoder = nn.Sequential(
            nn.Linear(latent_dim, input_dim),
            nn.Unflatten(1, (1, 28, 28))
        )

    def forward(self, X):
        # TODO: Implement the forward call
        latent = self.encoder(X)
        return self.decoder(latent), latent
```


For the training we can use the same function as for the convolutional autoencoder.

```
In [66]: n_epochs = 60
         lr = 1e-3
         # TODO: Train the PCA autoencoder model with a latent dimension of 10


         latent_dim = 10
         pca_ae_model = PCA_AE(28*28, latent_dim)
         train_loss, test_loss = train_ae(pca_ae_model, train_loader, test_loader,
```

Epoch 1/60 - Train Loss: 0.6656, Test Loss: 0.5735  
Epoch 2/60 - Train Loss: 0.4680, Test Loss: 0.3745  
Epoch 3/60 - Train Loss: 0.3228, Test Loss: 0.2815  
Epoch 4/60 - Train Loss: 0.2639, Test Loss: 0.2404  
Epoch 5/60 - Train Loss: 0.2334, Test Loss: 0.2141  
Epoch 6/60 - Train Loss: 0.2115, Test Loss: 0.1945  
Epoch 7/60 - Train Loss: 0.1939, Test Loss: 0.1805  
Epoch 8/60 - Train Loss: 0.1832, Test Loss: 0.1713  
Epoch 9/60 - Train Loss: 0.1748, Test Loss: 0.1647  
Epoch 10/60 - Train Loss: 0.1687, Test Loss: 0.1599  
Epoch 11/60 - Train Loss: 0.1647, Test Loss: 0.1560  
Epoch 12/60 - Train Loss: 0.1593, Test Loss: 0.1522  
Epoch 13/60 - Train Loss: 0.1551, Test Loss: 0.1488  
Epoch 14/60 - Train Loss: 0.1520, Test Loss: 0.1453  
Epoch 15/60 - Train Loss: 0.1478, Test Loss: 0.1418  
Epoch 16/60 - Train Loss: 0.1442, Test Loss: 0.1382  
Epoch 17/60 - Train Loss: 0.1403, Test Loss: 0.1345  
Epoch 18/60 - Train Loss: 0.1362, Test Loss: 0.1311  
Epoch 19/60 - Train Loss: 0.1328, Test Loss: 0.1279  
Epoch 20/60 - Train Loss: 0.1294, Test Loss: 0.1248  
Epoch 21/60 - Train Loss: 0.1268, Test Loss: 0.1223  
Epoch 22/60 - Train Loss: 0.1249, Test Loss: 0.1202  
Epoch 23/60 - Train Loss: 0.1220, Test Loss: 0.1184  
Epoch 24/60 - Train Loss: 0.1200, Test Loss: 0.1169  
Epoch 25/60 - Train Loss: 0.1189, Test Loss: 0.1154  
Epoch 26/60 - Train Loss: 0.1177, Test Loss: 0.1143  
Epoch 27/60 - Train Loss: 0.1168, Test Loss: 0.1134  
Epoch 28/60 - Train Loss: 0.1157, Test Loss: 0.1125  
Epoch 29/60 - Train Loss: 0.1146, Test Loss: 0.1118  
Epoch 30/60 - Train Loss: 0.1142, Test Loss: 0.1112  
Epoch 31/60 - Train Loss: 0.1135, Test Loss: 0.1105  
Epoch 32/60 - Train Loss: 0.1129, Test Loss: 0.1100  
Epoch 33/60 - Train Loss: 0.1124, Test Loss: 0.1095  
Epoch 34/60 - Train Loss: 0.1117, Test Loss: 0.1091  
Epoch 35/60 - Train Loss: 0.1114, Test Loss: 0.1087  
Epoch 36/60 - Train Loss: 0.1111, Test Loss: 0.1084  
Epoch 37/60 - Train Loss: 0.1106, Test Loss: 0.1079  
Epoch 38/60 - Train Loss: 0.1102, Test Loss: 0.1075  
Epoch 39/60 - Train Loss: 0.1093, Test Loss: 0.1071  
Epoch 40/60 - Train Loss: 0.1085, Test Loss: 0.1067  
Epoch 41/60 - Train Loss: 0.1085, Test Loss: 0.1063  
Epoch 42/60 - Train Loss: 0.1082, Test Loss: 0.1062  
Epoch 43/60 - Train Loss: 0.1078, Test Loss: 0.1054  
Epoch 44/60 - Train Loss: 0.1071, Test Loss: 0.1050  
Epoch 45/60 - Train Loss: 0.1072, Test Loss: 0.1045  
Epoch 46/60 - Train Loss: 0.1066, Test Loss: 0.1041  
Epoch 47/60 - Train Loss: 0.1059, Test Loss: 0.1037  
Epoch 48/60 - Train Loss: 0.1057, Test Loss: 0.1032  
Epoch 49/60 - Train Loss: 0.1052, Test Loss: 0.1027  
Epoch 50/60 - Train Loss: 0.1045, Test Loss: 0.1023  
Epoch 51/60 - Train Loss: 0.1043, Test Loss: 0.1019  
Epoch 52/60 - Train Loss: 0.1035, Test Loss: 0.1015  
Epoch 53/60 - Train Loss: 0.1031, Test Loss: 0.1011  
Epoch 54/60 - Train Loss: 0.1032, Test Loss: 0.1009  
Epoch 55/60 - Train Loss: 0.1028, Test Loss: 0.1005  
Epoch 56/60 - Train Loss: 0.1027, Test Loss: 0.1001  
Epoch 57/60 - Train Loss: 0.1023, Test Loss: 0.0999  
Epoch 58/60 - Train Loss: 0.1019, Test Loss: 0.0997  
Epoch 59/60 - Train Loss: 0.1018, Test Loss: 0.0996  
Epoch 60/60 - Train Loss: 0.1018, Test Loss: 0.0995

```
In [67]: # TODO: Visualize the training and test loss
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.title("Training Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.plot(train_loss, label=f"Latent Dim: {latent_dim}")
plt.legend()
plt.subplot(1, 2, 2)
plt.title("Test Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.plot(test_loss, label=f"Latent Dim: {latent_dim}")
plt.legend()
plt.tight_layout()
plt.show()
```



```
In [70]: # TODO: Visualize the reconstructed images and the original images
visualize_reconstruction(pca_ae_model, test_loader, device)
```



The results for the PCA reconstruction are worse compared to the CNN and the differences for the shape of the different classes are barely visible.

## Task 4

```
In [107... # TODO: Choose the convolutional autoencoder with latent dimension 3 and
conv_ae_model = models[latent_dims.index(3)]
conv_ae_model.eval()

latent_samples = []
```

```

latent_labels = []

with tc.no_grad():
    for X, Y in train_loader:
        X = X.to(device)
        output, latent = conv_ae_model(X)
        latent_samples.append(latent.cpu().numpy())
        latent_labels.append(Y.cpu().numpy())

        if len(latent_samples) * train_batch_size >= 800:
            break


latent_samples = np.concatenate(latent_samples)[:800]
latent_labels = np.concatenate(latent_labels)[:800]

```

```

In [110]: # TODO: Make a 3D scatter plot of the code colored by the labels
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(latent_samples[:, 0], latent_samples[:, 1], latent_samples[:, 2],
           c=latent_labels)
ax.set_xlabel('Latent Dimension 1')
ax.set_ylabel('Latent Dimension 2')
ax.set_zlabel('Latent Dimension 3')
plt.title('3D Scatter Plot of Latent Space')
plt.colorbar(ax.scatter(latent_samples[:, 0], latent_samples[:, 1], latent_labels),
            label='Labels')
plt.tight_layout()
plt.show()

```



As one can see, there happens a clustering into two clusters but the points are still

mixed together. But the clustering agrees with the underlying images, because classes 0 to 4 are all clothes and the shoes and bags (classes 5 to 9) are in the other cluster.